

Java VM

Terence Parr

partially derived with permission from Bill Venner's book,

"Inside the Java virtual machine"

<http://www.artima.com/insidejvm/ed2/>

Architecture

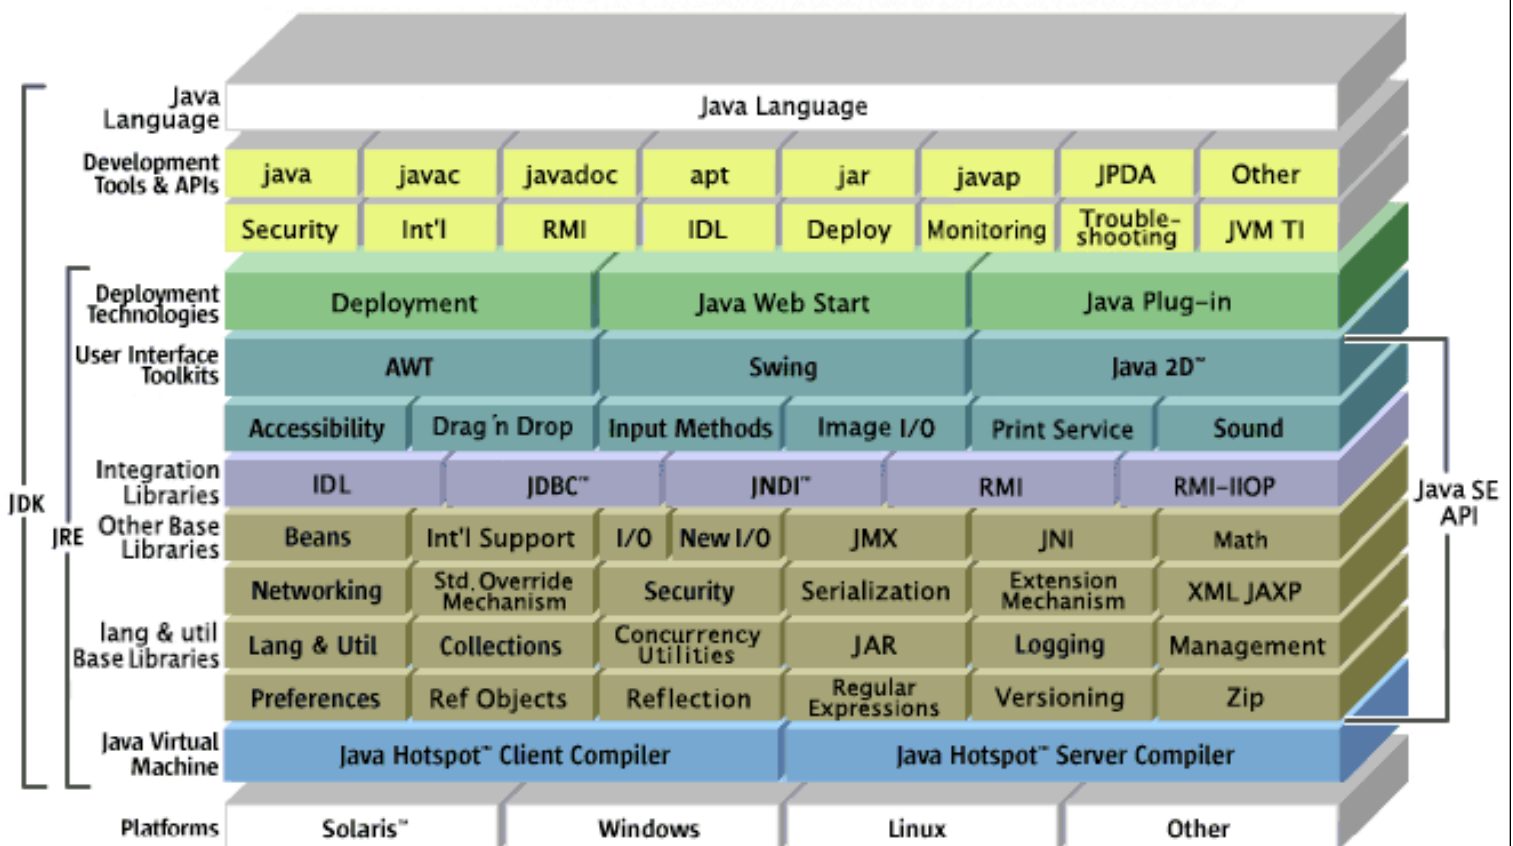
- Java has 4 key elements:
 - language
 - .class file format
 - same API across machines
 - virtual machine

Platform

- API/VM provide cross-platform execution
- VM isolates the program from x68, PPC, ...
- API mostly in Java with calls to native libraries via JNI

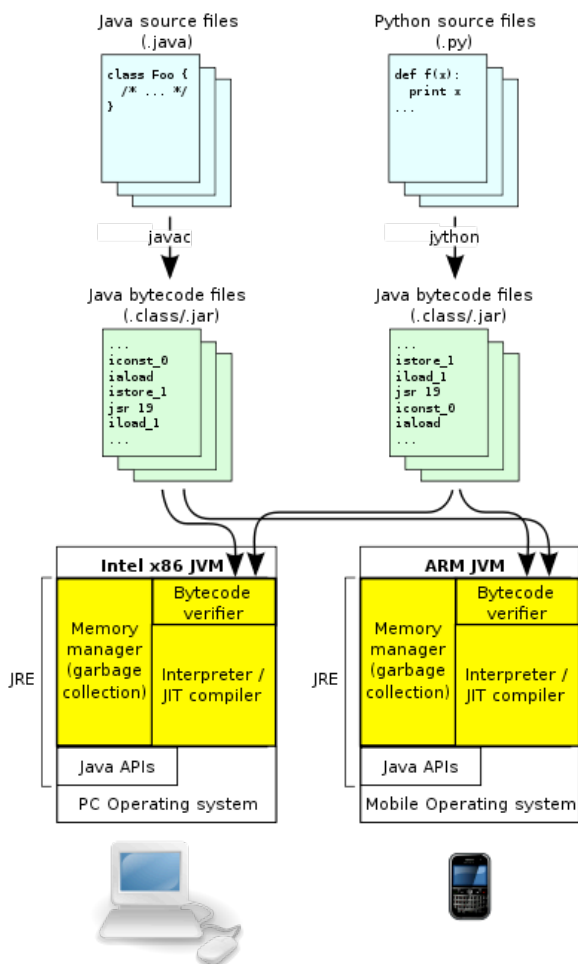
Platform II

Java™ Platform Standard Edition



<http://www.oracle.com/technetwork/java/whitepaper-135217.html>

Retargeting the JVM



Can run Clojure, Scala,
etc... on the Java VM

Program lifecycle

1. VM launches via “java MyClass”
2. ClassLoader loads java libraries and then MyClass on-demand (there is a boot loader)
3. Verification of bytecodes
4. Invoke main()
5. Create instances, GC tosses them out
6. Unload classes
7. VM exit when all user-level threads stop

Class loaders

- Finds .class files for a Class instance
- `java -verbose` to see where classes are loaded from
- Bootstrap loader is part of the VM (in C?)
- Bootstrap loads the core classes; can't redefine Object
- User-defined class loaders: can load classes from disk (our project) or by downloading classes from a website

Loading classes manually

```
loader = new URLClassLoader(  
    new URL[] { url1, url2 },  
    parentLoader);
```

```
final Class mainClass = loader.loadClass(className);  
final Method mainMethod =  
    mainClass.getDeclaredMethod(methodName);
```


Class loaders II

- Each class loader has a set of class definitions; own namespace
- Same class can be loaded into multiple loaders
- Classes are loaded on demand so linking actually occurs as the program executes, unlike C, which statically loads:
`cc -o myprog a.o b.o c.o`
- Care must be taken that CLASSPATH is the same on development and deployment computers
- Classes unloaded when the class loader is a candidate for GC

Verification

- JVM does not assume valid/safe compiler
- Very specific verification rules that sometimes disallow valid Java programs; e.g., when incident edges flow graph state have different stack sizes, verification fails even if it would be okay.
- Verifier verifies a lot of things including:
 - the stack does not grow forever (verifies size)
 - branches to valid locations (within same a method)
 - data is always initialized, type safe; cannot convert an integer to address
 - access modifiers such as private are respected (dynamic)
- Tries to prevent crashes, access to trusted code from on trusted applications

see also: <http://www.artima.com/insidejvm/ed2/lifetype3.html>

Class init

- Init direct superclass if hasn't been done
- Execute any static code blocks
- Initializing interface means executing static code block
- The initializer's poor class also initialize static fields to default values or specified values
- The initializer also creates static array objects with specified initialization values

Instance instantiation

- via: new, reflection, clone(), or deserializing object
- creates space on the heap for the instance variables:
 - this class and the superclass' fields
 - the bookkeeping variables (class ptr, lock)
- call <init>() to initialize object (if via “new” op)
- init:
 - call super.init
 - initialize instance variables
 - bytecode for constructor specifies in java

Runtime verification

- Check type casts
- No pointer arithmetic
- GC; no free(); no still pointers
- Array bounds checking
- Null dereference checks


Runtime security

`java.lang.SecurityManager`

- Remote classes in Applets execute within a restricted sandbox
- JVM tries to prevent malicious code from executing on the client
- Applets cannot
 - read/write to the local file system
 - make network connection to any host, except the host from which the applet came
 - create new process
 - load a new dynamic library and directly call native method
- Customizable

JIT, HotSpot

- JVM initially interprets bytecodes
- At a threshold, dynamic compiler translates to native code
- Key principle: vast majority of time spent in a minority of code
- Can re-optimize later if it's truly a hotspot
- Debugging complicates this; don't want to run in interpret only mode
- Some JVMs immediately compile code instead of initially interpreting (JIT)



JIT vs Dynamic compilation

- JIT not as good as dynamic compilation: consider that variables at static and JIT time are constants at runtime
- Why waste time compiling code that runs once or never?
- Compiling all code is an unnecessary waste of memory.
- By observing the program, we can gain useful profiling information such as:
 - how many times loops go around
 - caller/callee for in-lining (inlining reduces method call overhead and creates bigger blocks of code for the optimizer to chew on)
- Might have to de-optimize then reoptimize when new classes are loaded due to new methods coming available

Client compiler

- 3 phases:
 - by code to high-level intermediate representation (HIR) in static single assignment (SSA) form to enable more optimizations
 - platform specific “back-end” generates a lower level IR (LIR)
 - LIR to machine code with peephole optimization
- Focuses on local code quality, not global optimization (expensive)

Server compiler

- High-end fully optimizing compiler (uses SSA) including:
 - dead code elimination, range-check elimination
 - loop invariant hoisting, loop unrolling
 - global code motion (move code from main path to more control dependent patterns)
 - common subexpression elimination
 - constant propagation
 - global value numbering (tag var/expr w/same value with same number)
- register allocator is a global graph coloring allocator
- highly portable, relying on a machine description file to describe all aspects of the target hardware

Class file structure

<http://docs.oracle.com/javase/specs/jvms/se7/jvms7.pdf>

```
ClassFile {  
    u4          magic; // 0xCAFEBAE  
    u2          minor_version;  
    u2          major_version;  
    u2          constant_pool_count;  
    cp_info     constant_pool[constant_pool_count-1];  
    u2          access_flags;  
    u2          this_class; // index into constant pool  
    u2          super_class;  
    u2          interfaces_count;  
    u2          interfaces[interfaces_count];  
    u2          fields_count;  
    field_info  fields[fields_count];  
    u2          methods_count;  
    method_info methods[methods_count];  
    u2          attributes_count;  
    attribute_info attributes[attributes_count]; // filename, etc...  
}
```

Sample bytecodes

```
public void print();
```

```
LineNumberTable:
```

```
line 6: 0
```

```
line 7: 10
```

```
Code:
```

```
Stack=2, Locals=1, Args_size=1
```

```
0:  getstatic  #2; //Field java/lang/System.out:Ljava/io/PrintStream;
```

```
3:  aload_0
```

```
4:  getfield   #3; //Field width:I
```

```
7:  invokevirtual  #4; //Method java/io/PrintStream.println:(I)V
```

```
10: return
```

```
LineNumberTable:
```

```
line 6: 0
```


```
line 7: 10
```

```
public class Shape {  
}  
class Square extends Shape {  
    int width;  
    public void print() {  
        System.out.println(width);  
    }  
}
```

Constant pool

```
0:  getstatic  #2; //Field java/lang/System.out:Ljava/io/PrintStream;
3:  aload_0
4:  getfield   #3; //Field width:I
7:  invokevirtual #4; //Method java/io/PrintStream.println:(I)V
```

Const pool #3 is width field:



cp_info #5 is Class name:

Class name: [cp_info #22](#) <Square>

#22 is Utf8:

Length of byte array: 6
Length of string: 6
String: Square

#19 is NameAndType:

Name: [cp_info #7](#) <width>
Descriptor: [cp_info #8](#) <I>

Method info

Name: [cp_info #13](#) <print>
Descriptor: [cp_info #10](#) <OV>
Access flags: 0x0001 [public]

Generic info:

Attribute name index: [cp_info #12](#)
Attribute length: 10

Specific info:

Nr.	start_pc	line_number
0	0	6
1	10	7

jasmin assembler

- <http://jasmin.sourceforge.net/>
- Assembles bytecodes (text) to .class files

```
.method public static main([Ljava/lang/String;)V
    .limit stack 2    ; System.out and "hi mom" on stack
    .limit locals 1   ; the arg

    getstatic        java/lang/System/out Ljava/io/PrintStream;
    ldc              "Hi mom"
    invokevirtual    java/io/PrintStream/println(Ljava/lang/String;)V
    return
.end method
```

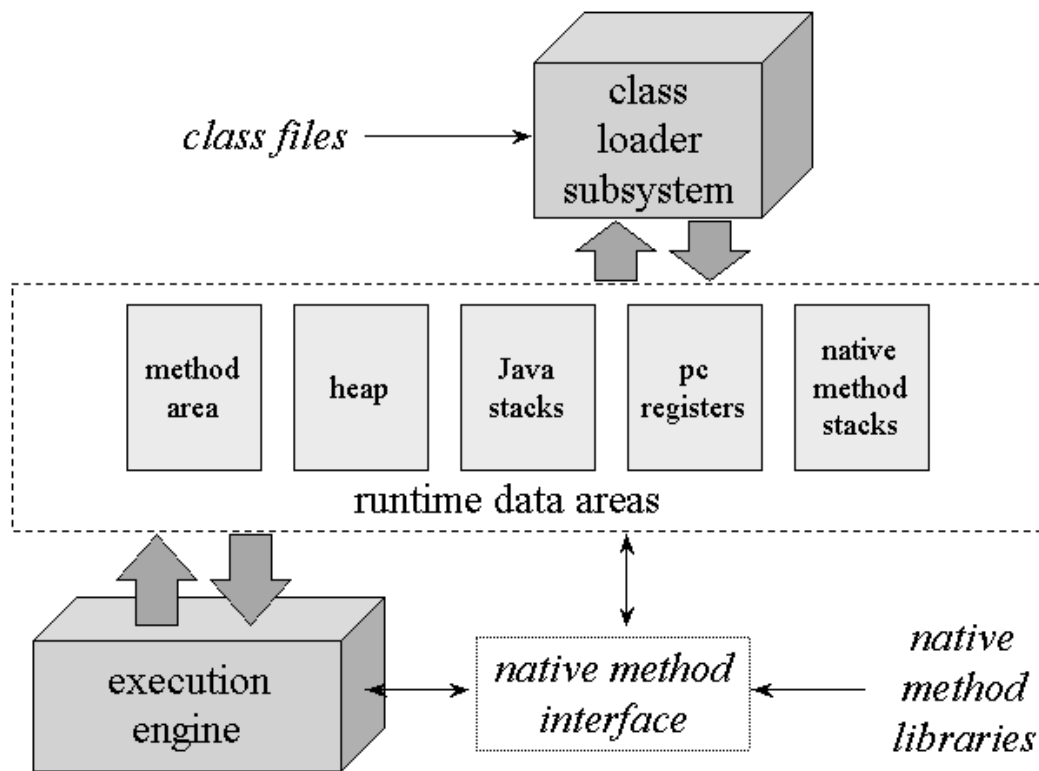
ASM

- <http://asm.ow2.org/>
- Visitor-based bytecode library
- We'll use this for a bytecode project

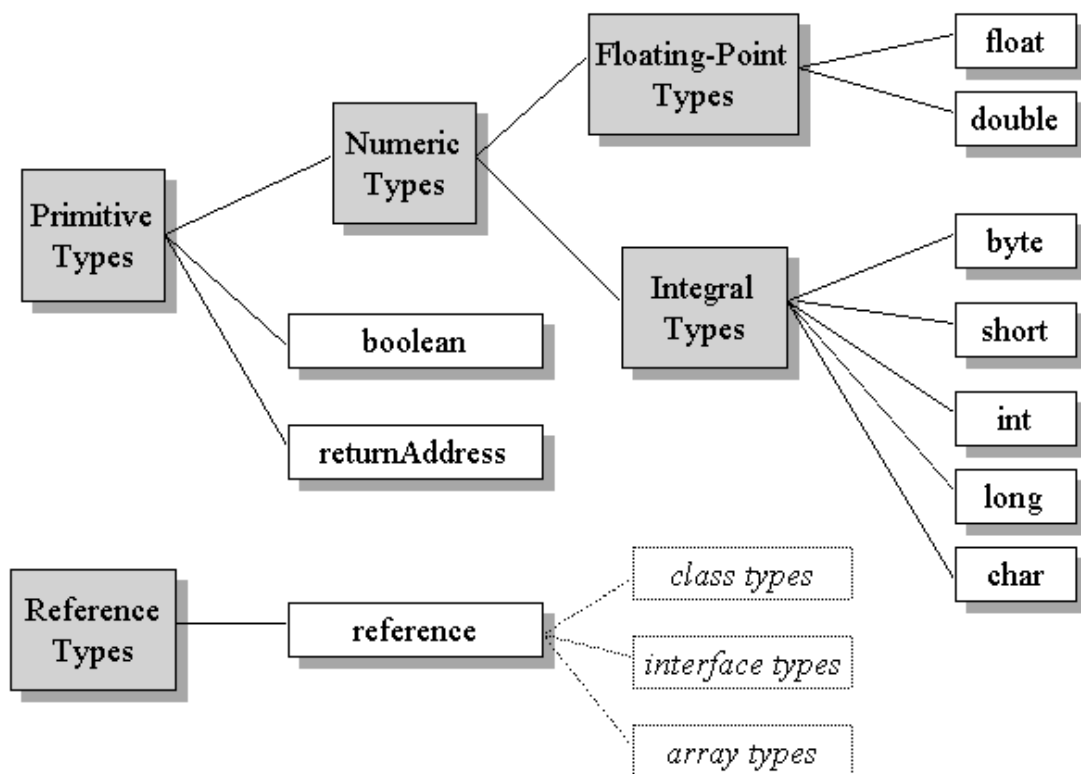
```
ClassReader cr = new ClassReader(is);
ClassWriter cw = new ClassWriter(0);
ClassVisitor cv = new MyVisitor(cw);
cr.accept(cv, 0); // visit
byte[] b = cw.toByteArray();
FileOutputStream fos = new FileOutputStream(...);
fos.write(b);
fos.close();
```


JVM execution engine

Architecture



Data types




VM elements

- Java stack (call stack) of method invocation frames
- A stack frame state of one method invocation; has:
 - local variables (size determined statically), parameters become local variables during method invocation
 - operand stack (32 bits, max size determined statically)
 - return values, return address, parent frame
 - other runtime information needed to perform dynamic linking such as pointer to the constant pool
- Current frame/method pointer
- Each thread gets its own java stack, pc, native code stack

Example method call

```
void foo() {  
    int x = bar(9,10);  
}  
int bar(int x, int y) {  
    return x+y;  
}
```



Heap

- All non-primitive types are allocated in the heap (via new)
- Primitive types can be allocated on the stack like C (so could objects in some cases)
- The heap is managed by the garbage collector
- No way to explicitly free something
- Objects live as long as there is a path to it from outside the heap

Instruction set

- One byte opcodes == Bytecodes
- 0 or more operands after bytecode (or encoded within the byte code)
- Some operands are on the operand stack
- Typed by first letter:
iload, lload, fload, dload, aload
- Compiler must encode byte, short sometimes (sign extend to int), boolean/char 0-extend
 - Only push,aload,astore,i2b,i2s take byte/short

<http://docs.oracle.com/javase/specs/jvms/se7/jvms7.pdf>

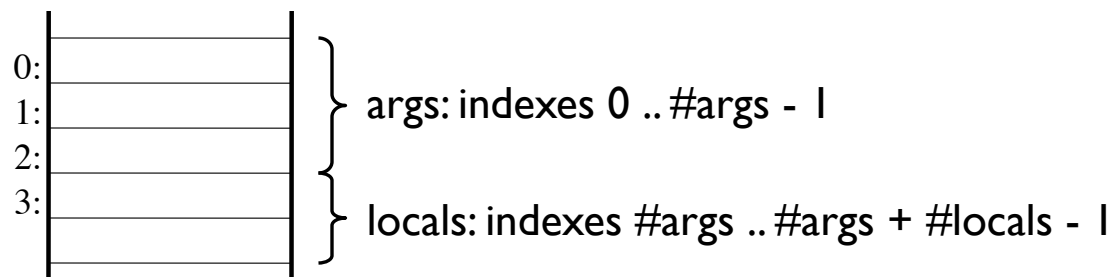
Load and store

iload forms:

```
iload_0  
iload_1  
iload_2  
iload_3  
iload n  
wide iload n
```

- Load local variable onto stack: `iload`, `iload_n`, `aload`, ...
- Store from operand stack into local: `istore`, `istore_n`, `astore`, ...
- *wide* modifier gives access to a wider index

Arguments, locals inside stack frame



http://cs.ua.edu/434/Classes/20_JVM.ppt

Arithmetic instructions

Arithmetic

add: iadd, ladd, fadd, dadd

subtract: isub, lsub, fsub, dsub

multiply: imul, lmul, fmul, dmul
etc.

Conversion

i2l, i2f, i2d,

l2f, l2d, f2d,

f2i, d2i, ...

http://cs.ua.edu/434/Classes/20_JVM.ppt

More...

Operand stack manipulation

`pop, pop2, dup, dup2, swap, ...`

Control transfer

Unconditional: `goto, jsr, ret, ...`

Conditional: `ifeq, iflt, ifgt, if_icmpeq, ...`

http://cs.ua.edu/434/Classes/20_JVM.ppt

Example loop

```
for (int i = 0; i < 100; i++) { }
```

```
0   iconst_0           // Push int constant 0
1   istore_1           // Store into local variable 1 (i=0)
2   goto 8             // First time through don't increment
5   iinc 1 1           // Increment local variable 1 by 1 (i++)
8   iload_1            // Push local variable 1 (i)
9   bipush 100         // Push int constant 100
11  if_icmplt 5         // Compare and loop if i<100
14  return             // Return void when done
```

<http://docs.oracle.com/javase/specs/jvms/se7/jvms7.pdf>

Example numeric value in constant pool

<code>int i = 100;</code>	<code>0 bipush 100</code>
<code>int j = 1000000;</code>	<code>2 istore_1</code>
<code>long l1 = 1;</code>	<code>3 ldc #1</code>
<code>long l2 = 0xffffffff;</code>	<code>5 istore_2</code>
<code>double d = 2.2;</code>	<code>6 lconst_1</code>
	<code>7 lstore_3</code>
	<code>8 ldc2_w #6</code>
	<code>11 lstore 5</code>
	<code>13 ldc2_w #8</code>
	<code>16 dstore 7</code>

<http://docs.oracle.com/javase/specs/jvms/se7/jvms7.pdf>

Non-local memory access

accessing locals and arguments: `load` and `store` instructions

accessing fields in objects: `getfield`, `putfield`

accessing static fields: `getstatic`, `putstatic`

Note: Static fields are a lot like global variables.

They are allocated in the “method area” where also code for methods and representations for classes (including method tables) are stored.

Note: `getfield` and `putfield` access memory in the heap.

http://cs.ua.edu/434/Classes/20_JVM.ppt

Heap Memory Allocation

Create new class instance (object):

`new`

Create new array:

`newarray`: for creating arrays of primitive types.

`anewarray`, `multianewarray`: for arrays of reference types.

http://cs.ua.edu/434/Classes/20_JVM.ppt

Calling methods

Method invocation:

`invokevirtual`: usual instruction for calling a method on an object.

`invokeinterface`: same as `invokevirtual`, but used when the called method is declared in an interface (requires a different kind of method lookup)

`invokespecial`: for calling things such as constructors, which are not dynamically dispatched (this instruction is also known as `invokenonvirtual`).

`invokestatic`: for calling methods that have the “static” modifier (these methods are sent to a class, not to an object).

Returning from methods:

`return`, `ireturn`, `lreturn`, `areturn`, `freturn`, ...

http://cs.ua.edu/434/Classes/20_JVM.ppt